NeuroMPS

A provenance-tracked MEA data asset and a calibrated two-compartment digital twin

Team	NeuroMPS
Category	AI4S Open Innovation: AI + Organ-on-a-Chip (5th Pazhou Algorithm Competition)
Repository	`https://github.com/Vodolay/AI4S`
Commit	`a5a24ca (working tree dirty)`
Built	2026-09-22

Research decision-support only. No efficacy, safety, diagnostic or clinical claim is made.

Contents

1. Summary **2**

2. Problem and scope **3**

3. The data standard **3**
3.1 Record model and data levels . . . 3
3.2 Provenance and conformance . . . 4
3.3 Quality control . . . 4
3.4 Feature set . . . 4

4. The data asset **5**
4.1 A rejection, as a worked example of the license discipline . . . 5
4.2 A licence conflict, resolved rather than avoided . . . 5

5. Methods **6**
5.1 Cross-validation and resampling . . . 6
5.2 Tasks . . . 6

6. Results: data and models **6**
6.1 Task 1 — mechanism classification . . . 6
6.2 Task 2 — external detection and its negative control . . . 7
6.3 Species transfer (rat → human) . . . 9
6.4 Task 3 — platform contrast . . . 9
6.5 Which features survive a species change . . . 10
6.6 Ablations: the feature set is mostly one feature . . . 10
6.7 Compartment asymmetry in real devices . . . 11

7. Results: the digital twin **12**
7.1 Model . . . 12
7.2 The measure the twin exists to use is confounded . . . 12
7.3 Calibration . . . 13
7.4 Drug knobs . . . 13
7.5 Simulation-based inference, and what it cannot do . . . 14
7.6 Directionality: what it takes to see asymmetry . . . 15
7.7 The same question on real devices, and a null . . . 16

8. Limitations **18**

9. Reproducibility **19**
9.1 Where the time goes . . . 19
9.2 What the repository enforces, rather than documents . . . 20

10. The demo **20**

11. AI assistance disclosure **20**

12. Sources and licenses **21**

Appendix A. Where each result lives **21**

Every number below carries an HTML comment naming the `outputs/metrics.json` key it came from. `python -m neuromps.report.verify` re-reads each citation and

fails if the document and the metrics disagree, so a figure that goes stale when a stage is regenerated cannot quietly survive in the prose. The decision record behind these results, including what was tried and abandoned, is `report/LAB_NOTEBOOK.md`.

---

1. Summary

Two layers were built:

(A) A standardized, provenance-tracked data asset of public neural microelectrode-array (MEA) recordings: 21,380 records from five sources, each record carrying its source URL, license, SHA256 and download time, validated against a JSON Schema at a conformance level that requires every provenance hash to match the download manifest.

(B) A Brian2 spiking-network digital twin of an asymmetric two-compartment neural chip, calibrated against the asset, with simulation-based inference (`sbi`) attempting to recover connectivity and drug effects from MEA summary features.

The results that matter are as much negative as positive, and the report leads with both:

1. **Rate and burst-fraction measures survive a change of species; synchrony and network-burst measures do not** (§6.3, §6.5, §6.6, §7.3). **Across platforms none of them is safe** — mean firing rate itself reverses sign (§6.4), which is the sharper limitation of the two.
2. **Coupling asymmetry tracks relative firing rate, not pathway direction.** At fixed wiring, changing one chamber's drive moves the measure several times further than changing the wiring does (§7.2).
3. **The twin's posterior does not recover drug-knob direction from real recordings.** The inferred GABA-A scale shows no bicuculline dose trend, and untreated chips already infer 0.3781, so the parameter is absorbing model mismatch rather than tracking the compound (§7.5).
4. **Direction is recoverable from simulated features once firing rate is residualised out** — AUC 0.9359, 95% CI [0.893, 0.9739], from a 10 s recording on 24 electrodes per compartment — but as a contrast between two well-spaced wiring classes, not as a measurement of a continuous ratio (§7.6). A second arm, which adjusted drive to match rates, reaches a higher AUC at fewer electrodes and **fails its own validity check**; it is reported as inconclusive rather than as the result.
5. **That result does not reproduce on real devices.** On 26 chips whose microchannel geometry was fabricated rather than assumed, propagation asymmetry does not separate straight channels from three asymmetric designs (Kruskal $p = 0.64$), and the statistic is confounded with how many events a chip produced (§7.7). The sample is too small to call that a refutation of the simulated result, and too directly aimed to omit.

The honest summary of the twin is therefore: it reproduces baseline features and most drug directions, it can tell two wiring regimes apart in simulation, it cannot yet read pharmacology off a real chip, and the one real-data test of its directionality claim is null.

2. Problem and scope

Microphysiological systems produce electrophysiology that is rich but hard to compare: different platforms, electrode counts, durations, species and feature definitions. Two consequences follow for anyone trying to use public MEA data for compound screening.

Data are not comparable by default. Feature names collide while definitions differ — "synchrony" from a spike-time tiling coefficient is not the Axion synchrony index — and provenance is usually lost by the time data reach an analysis.

A model of the chip is not a model of the measurement. A spiking network that reproduces a culture's firing does not necessarily reproduce what 48 electrodes detect from it.

This entry addresses both: a schema and feature standard with enforced provenance, and a digital twin that is scored through the same feature extractor as the real recordings.

Out of scope. No clinical, diagnostic, efficacy or safety claim is made or implied. Compound labels are mechanism classes drawn from the source publications, not indications.

3. The data standard

3.1 Record model and data levels

Records validate against `schema/neuromps_schema.json` (JSON Schema 2020-12). The central decision is that records carry an explicit **data level**, and that records only compare within a level:

level	meaning	records
`spike_times`	per-electrode spike trains; our feature vector is computable	1,005
`fitted_curve`	published concentration-response fits only	20,375
`per_well_features`	vendor feature tables without spike times	—

This distinction is load-bearing. Most of the asset by count is ToxCast curve fits, from which our nine-feature vector cannot be recomputed; only 897 records support the feature pipeline. A schema that hid that behind a single "features" field would invite exactly the wrong comparison.

3.2 Provenance and conformance

Conformance **level 2** requires that every record's provenance hash matches `data/raw/MANIFEST.json`, and the build refuses to emit an asset otherwise. Each fetcher records URL, license, SHA256, byte count and download time, and re-checks the source's advertised license at fetch time so a silently-relicensed upstream blocks rather than propagates.

3.3 Quality control

QC is recorded, never applied silently. Each record carries flags and a single `qc.usable` verdict, and the flags are informational: a flagged record is still in the asset, with the reason attached, so a downstream analysis can make its own decision.

flag	records	what it means
`incomplete_dose_series`	272	the compound's dose series is missing points
`too_few_active_electrodes`	184	under 10 electrodes cleared 5 spikes/min
`replicate_divergence`	161	one replicate departs from its group
`outlier_feature_value`	161	a feature sits far outside its source's distribution
`excluded_well_upstream`	90	the source's own log excluded the well
`noisy_electrodes_present`	11	the source marked electrodes as noisy

Two decisions here were driven by defects found in the data rather than by design. The replicate-divergence test originally used a robust z-score, which collapses when more than half a replicate group shares a value — common in sparse recordings, where the median absolute deviation is exactly zero; it now falls back to a mean absolute deviation and names the diverging feature. And the 90 wells excluded by an upstream log were at first missing from the asset altogether, because records were built from the feature tables alone; they are now built from the union of the log and the spike files, so an exclusion is represented rather than absent.

3.4 Feature set

Nine features form the core vector, computed by a single extractor (`features/mea.py`) that the twin must also call:

feature	definition as computed
`mfr_hz`	spikes per second per active electrode
`burst_rate_per_min`	bursts per minute per active electrode
`burst_duration_s`	mean burst duration
`spikes_per_burst`	mean spikes in a burst
`pct_spikes_in_bursts`	share of spikes falling inside bursts
`network_burst_rate_per_min`	network bursts per minute, chance-corrected participation
`synchrony_index`	spike-time tiling coefficient, mean over electrode pairs

feature	definition as computed
`isi_cv`	coefficient of variation of inter-spike intervals
`active_electrodes`	electrodes clearing the activity threshold

Thresholds are fixed and recorded per record: an electrode is active at ≥5 spikes/min, a well is usable at ≥10 active electrodes, a burst is ≥5 spikes with ISI ≤0.1 s, and the tiling coefficient uses $\Delta t = 0.05$ s.

The network-burst threshold deserves its own note, because the obvious implementation is wrong. A fixed participation threshold — n electrodes firing in the same window — counts coincidences that independent electrodes produce anyway; at 5 Hz on 16 electrodes that is on the order of fifty spurious "network bursts" per minute, which would dominate the measure. The threshold is therefore set per recording so that the expected number of chance detections stays below 0.05, from the recording's own electrode count and firing rates. This matters beyond bookkeeping: because the correction scales with electrode count, it is a candidate explanation for part of the cross-platform disagreement reported in §6.4, and that remains untested.

---

4. The data asset

source	records	level	species	license
`toxcast_mea`	20,375	fitted_curve	rat	CC0-1.0
`screenneuropharm_bic`	580	spike_times	rat	MIT
`kapucu2022`	288	spike_times	human, rat	CC-BY-4.0
`sharf2022`	29	spike_times	human	CC0-1.0
`mateus2024`	108	spike_times	rat	CC-BY-4.0

Totals: 21,380 records, of which 20,924 are marked usable and 456 are not. By species, 21,207 records are rat and 173 are human.

The species imbalance is the asset's main structural weakness and is stated wherever human results appear: every human-species conclusion rests on 173 records from two sources.

4.1 A rejection, as a worked example of the license discipline

Tukker 2020 (rat and hiPSC MEA pharmacology) is a good fit scientifically and is **excluded**. It is CC BY-NC 4.0, and the competition rules assign joint ownership of submissions rather than preserving an entrant's own license, so the non-commercial condition cannot be guaranteed downstream. It is recorded as `skipped` in the manifest with that reason rather than dropped silently. The cost is real: without it, the rat→human transfer study rests on one source.

4.2 A licence conflict, resolved rather than avoided

Zenodo 14525182 (Mateus et al. 2024) is the closest published analogue of the chip the twin models: hippocampal cultures on straight against asymmetric microchan-

nels, 57 spike-data recordings, with the geometry fabricated into the device. It was **rejected on 2026-09-16** because the deposit contradicts itself — the Zenodo metadata licence field says `cc-by-4.0` and the record is open access, while the README inside says, in capitals, CC BY-NC-ND. Under NoDerivatives, publishing derived features would not have been permitted, which is the whole of what this project does with a recording.

On 2026-09-21 the corresponding author confirmed by email that the metadata governs and the README line was an unintentional default. The source is included on that basis, and the reasoning is recorded where a reader can check it: both halves of the contradiction are public, the fetcher blocks if the metadata licence ever moves, and it reads the README on every fetch and writes whatever licence claim it finds into `data/raw/MANIFEST.json`. **The resolution itself is private correspondence, which a reader cannot verify**; `data/SOURCES.md` says so in those words. This is the one source in the asset whose licence rests on anything other than a public record.

Its 108 records are in the asset: two chambers per recording, split at the microchannel band recovered from the spike times, with features from the same extractor as every other source. They are analysed in §7.7.

---

5. Methods

5.1 Cross-validation and resampling

Every model is evaluated with **grouped** cross-validation, grouping by culture/well so replicates of one culture never span the train/test boundary, and every interval is a **cluster bootstrap over those groups**, never over rows — with 1,000 resamples in a full run and 200 in `--quick`, the count recorded alongside each estimate. Bootstrapping rows here would report intervals roughly a factor of two too narrow, because rows within a culture are not independent.

5.2 Tasks

- **Task 1 — mechanism classification** (within platform): predict mechanism class from the nine-feature delta vector of a treated/baseline pair.
- **Task 2 — external detection**: apply the Task 1 model to a source it never saw, with a negative control of the opposite mechanism.
- **Task 3 — platform contrast**: compare the same mechanism measured on two platforms.
- **Transfer**: rat → human, with and without domain adaptation.

---

6. Results: data and models

6.1 Task 1 — mechanism classification

Six classes, 73 treated/baseline pairs across 50 culture groups, 4-fold grouped CV.

model	balanced accuracy	95% CI
logistic regression	0.537	[0.4216, 0.6845]
gradient boosting	0.398	[0.2762, 0.5178]
most-frequent baseline	0.167	—

Chance is 0.167. Logistic regression beats chance with an interval that excludes it. Its interval and gradient boosting's overlap heavily, so the gap between the two models is within noise; with 73 pairs the simpler model is the one to prefer, not the one shown to be better.

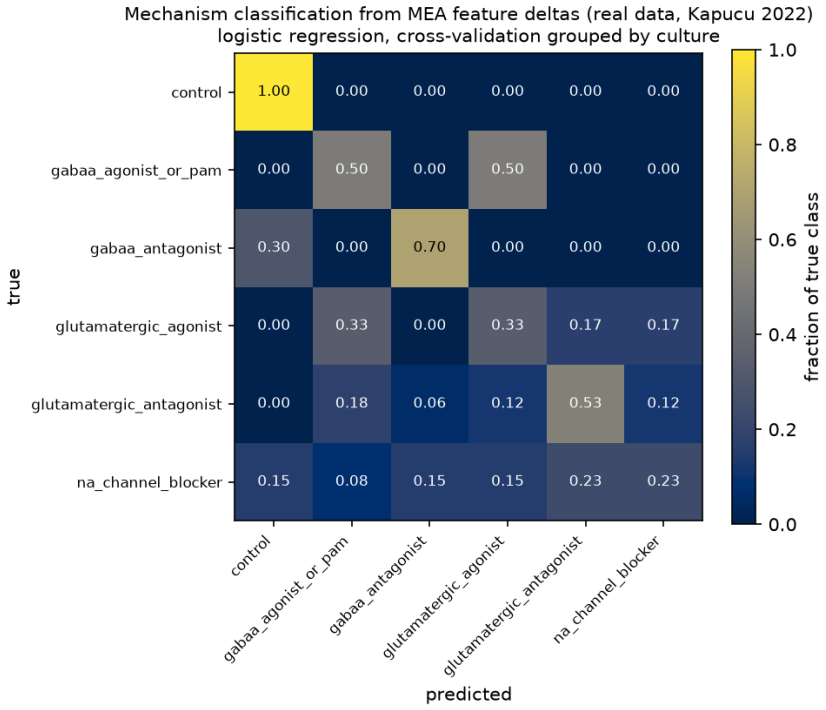


Figure 1: Task 1 confusion matrix, REAL data: out-of-fold predictions from grouped cross-validation over 73 treated/baseline pairs. Errors concentrate between the two GABA-A classes.

Calibration. Expected calibration error on out-of-fold predictions is 0.1114, 95% CI [0.0773, 0.2557] (`outputs/figures/task1_reliability_real.png`). The upper end of that interval is large: predicted probabilities should be read as ordering, not as calibrated risk — and not even as ordering above 0.8 confidence, where observed accuracy falls rather than rises, on a handful of predictions.

Conformal sets. At a 95% target the empirical coverage is 0.918 with a mean set size of 4.562 classes out of six — honest coverage at the cost of sets too large to act on. That is the correct reading of 73 calibration points, not a defect to tune away.

6.2 Task 2 — external detection and its negative control

Applied to bicuculline recordings from a source the model never saw, the GABA-A antagonist class is detected in 0.366 of 497 records, and the detection does **not** track dose (Spearman 0.014, p = 0.751).

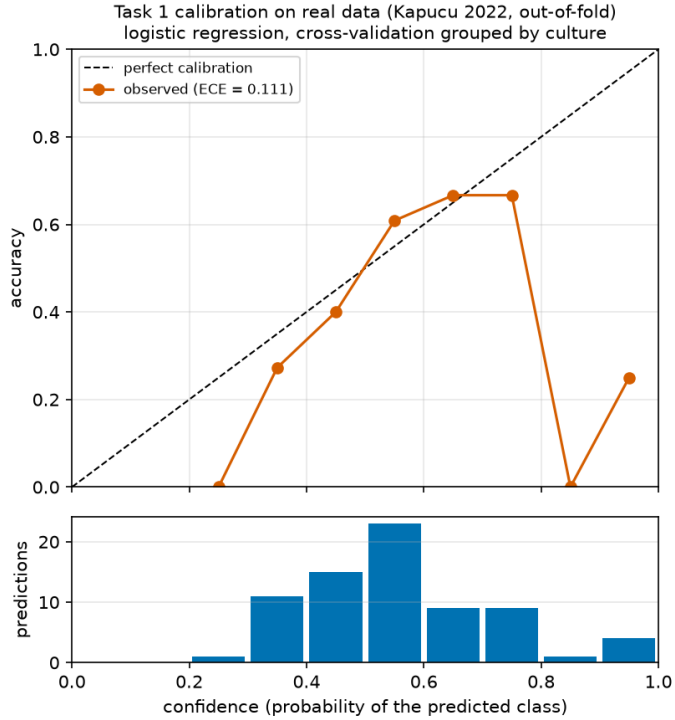


Figure 2: Reliability diagram, REAL data: predicted probability against observed frequency for the mechanism classifier, out-of-fold. Bin counts are small, which is why the calibration interval is wide.

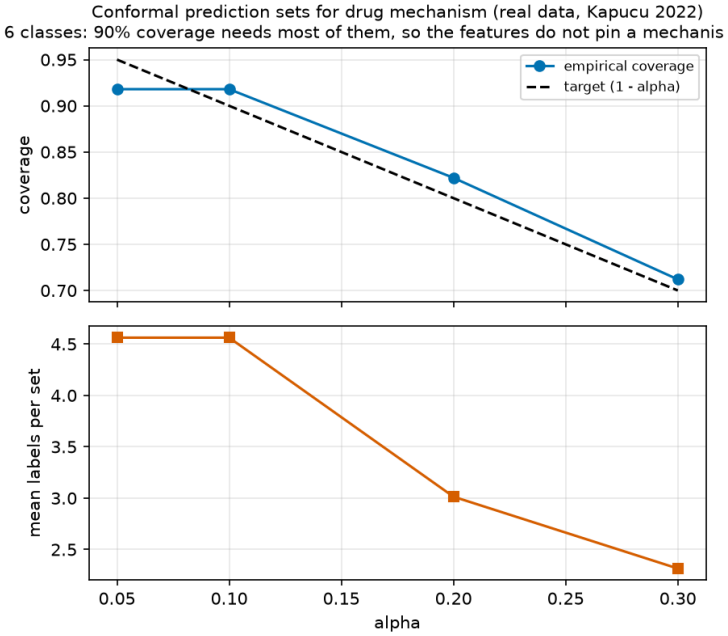


Figure 3: Conformal set size against target coverage, REAL data. Honest coverage costs sets of four to five classes out of six.

The negative control behaves as it should: diazepam, a positive modulator rather than an antagonist, is called GABA-A antagonist in only 0.077 of its records.

This is a weak result reported as a weak result. A detection rate near a third with no dose relationship is not a screening tool; it is evidence that the mechanism signal partially survives a platform change.

6.3 Species transfer (rat → human)

setting	balanced accuracy	95% CI
human, within-species	0.646	—
rat, within-species	0.543	—
rat → human, no adaptation	0.55	[0.4, 0.6]
rat → human, CORAL	0.442	[0.28, 0.597]
rat → human, per-species standardisation	0.408	[0.246, 0.56]

Chance is 0.2. Transfer works above chance (permutation p = 0.0002), and **both adaptation methods made it worse**. With 37 training and 32 test records, CORAL is estimating a covariance it cannot support. Reporting adaptation as a negative result is the useful outcome here.

6.4 Task 3 — platform contrast

The same mechanism at ≥30 µM measured on two platforms: gabazine on 16-electrode Axion wells, n = 10, against bicuculline on multi-electrode arrays, n = 167. Five features agree in sign and four disagree.

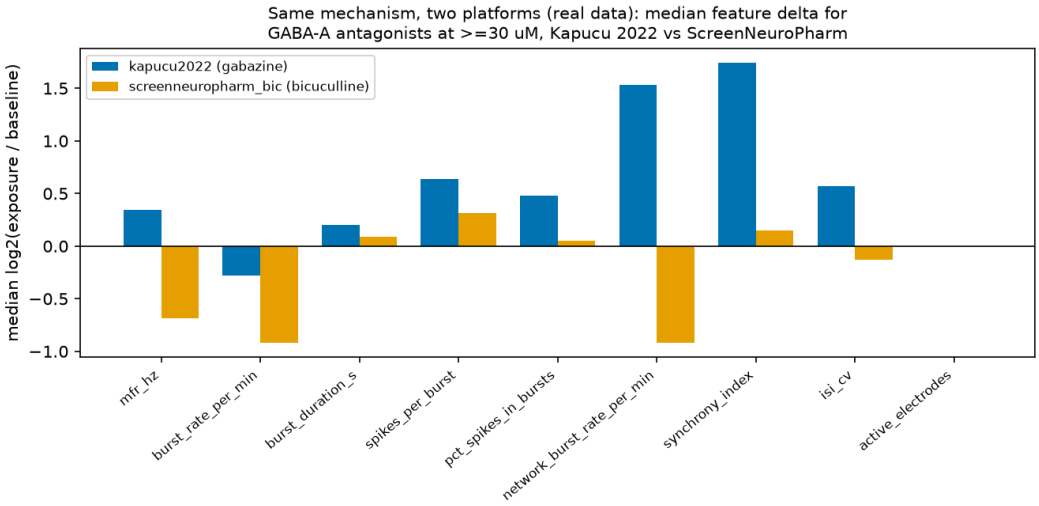


Figure 4: The same mechanism measured on two platforms, REAL data: per-feature log2 change for GABA-A antagonists. Mean firing rate reverses sign across platforms; synchrony agrees.

The disagreement does not fall where the species analysis puts it. `mfr_hz` reverses, +0.348 against −0.684, while `synchrony_index` agrees, +1.743 against +0.148. That is the opposite of the pattern in §6.5 and §6.6, and it is the reason this report separates the two claims: **the portable-feature result is about species, and does not carry across platforms.**

Three cautions on how far this goes. The contrast confounds compound (gabazine against bicuculline) with electrode count, recording duration and culture age, so it bounds portability rather than explaining it. The two arms are ten wells against 167 records. And `active_electrodes` is counted as a disagreement while its median delta is exactly 0.000 on both platforms — a zero has no sign, so the substantive count is three disagreements, not four.

6.5 Which features survive a species change

The most shifted features between species are `synchrony_index`, `isi_cv` and `network_burst_rate_per_min`; the most invariant are `mfr_hz`, `pct_spikes_in_bursts` and `burst_rate_per_min` (`outputs/figures/species_shift.png`). The synchrony shift is large — a difference of −4.179 in log2 units — though its interval [−5.7435, 0.1868] reaches zero, so the ordering is better supported than the magnitude. The same hedge applies to `isi_cv`, whose interval [−1.6738, 0.15] also reaches zero; only `network_burst_rate_per_min` has an interval that excludes it, [−9.7879, −0.3389]. The ranking is by standardised difference, so "most invariant" does not mean small: `burst_rate_per_min` is named invariant while carrying a median difference of 1.04 log2 units, a two-fold shift whose interval also spans zero.

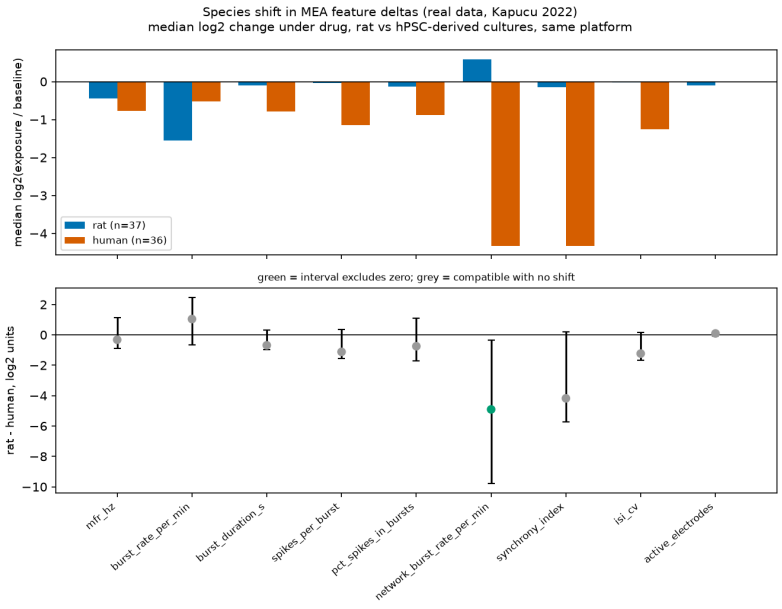


Figure 5: Per-feature species shift, REAL data: rat against human baseline distributions, with cluster-bootstrap intervals. Rate and burst-fraction features move least; synchrony moves most.

6.6 Ablations: the feature set is mostly one feature

feature subset	balanced accuracy	95% CI
all nine	0.514	[0.3983, 0.6619]
rate only	0.53	[0.3964, 0.6235]
species-invariant three	0.516	[0.371, 0.6295]
synchrony only	0.455	[0.3554, 0.586]
network measures only	0.417	[0.2983, 0.5573]

Two rate features match all nine. Leave-one-out agrees: dropping `mfr_hz` changes balanced accuracy by −0.05, while dropping `burst_rate_per_min` changes it by 0.0.

6.7 Compartment asymmetry in real devices

Across 16 two-chamber devices, the statistic that separates device types is coupling **lag**, not asymmetry. Median absolute lag is 11.25 ms for cortical–hippocampal devices against 5.0 ms and 5.0 ms for the two same-tissue types (Kruskal p = 0.0082; Holm-adjusted 0.0413 and 0.0257 against the two homogeneous types).

The internal control behaves. The two same-tissue types are indistinguishable from each other, Holm 0.4785, which is what a real device-type effect should look like and what a pooled heterogeneous-against- homogeneous split cannot show.

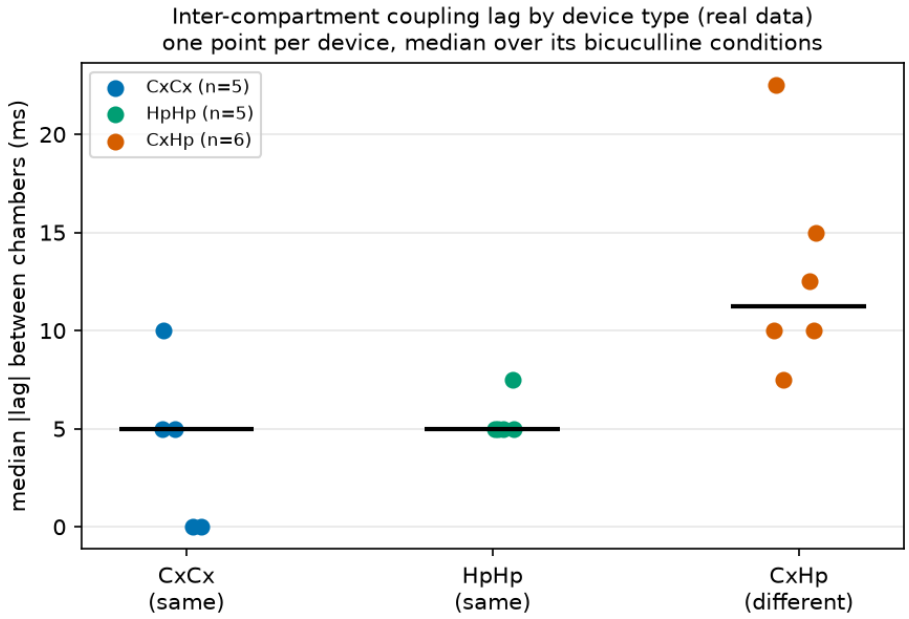


Figure 6: Inter-compartment coupling lag by device type, REAL data, one value per device: five CxCx, five HpHp, six CxHp. The two same-tissue types are indistinguishable from each other; the heterogeneous type sits apart.

Two things this section does **not** claim. On absolute asymmetry rather than lag, the same per-type comparison does not survive correction, Holm 0.0519, so lag is the claim and asymmetry is not. And a pooled heterogeneous-against-homogeneous test of asymmetry does cross the line, p = 0.042; it is kept in `metrics.json` and is deliberately not the headline here, because pooling CxCx with HpHp is an artefact — the two homogeneous types differ from each other on peak correlation, so pooling them manufactures a contrast. An earlier draft of this report led with that pooled test; the lab notebook had already recorded it as abandoned, and this section now reports the analysis that replaced it.

With five to six devices per group, this supports an ordering of device types, not an effect size, and §7.2 gives a specific reason to distrust the asymmetry measure in particular.

---

7. Results: the digital twin

7.1 Model

Two compartments of conductance-based leaky integrate-and-fire neurons with spike-frequency adaptation, 80/20 excitatory/inhibitory, recurrent intra-compartment connectivity and **directional** inter-compartment projections whose forward/backward ratio is the parameter of interest. Poisson background drive. Connectivity is parameterised by **in-degree rather than probability**, so a calibration done at one network size still holds at another.

The parameters carried by the network, and how each was fixed:

parameter	role	how it was set
`intra_in_degree`	recurrent inputs per neuron	calibrated, searched over 10-35
`inter_in_degree`	cross-compartment inputs per neuron	inferred, prior 4-16
`forward_ratio`	forward against backward pathway strength	inferred, prior 1-16 in log space
`excitatory_weight`, `inhibitory_weight`	synaptic conductance steps	calibrated
`background_rate_hz`, `background_weight`	Poisson drive	calibrated; drive to the target chamber is a separate scale
`adaptation_increment`, `adaptation_tau_ms`	spike-frequency adaptation	calibrated
`threshold_spread`	per-neuron excitability heterogeneity	calibrated, searched over 0-0.3
`detection_radius_um`	how far an electrode hears	calibrated, searched over 100-300 µm
`gabaa_scale`, `ampa_scale`, `threshold_scale`	drug knobs, 1.0 is untreated	inferred

Identical neurons synchronise almost perfectly, which no real culture does, so `threshold_spread` exists to stop the twin being trivially more synchronous than any recording. `detection_radius_um` is a property of the virtual electrode rather than of the culture, and is called out as such in the calibration record so it is not read as a biological parameter.

Neurons are placed on a virtual electrode grid and only spikes within a detection radius reach an electrode; the resulting per-electrode spike trains go through the same `features/mea.py` as real recordings, and — after a defect found in HD16 — through the same coupling measure on the same per-electrode input.

7.2 The measure the twin exists to use is confounded

Coupling asymmetry was expected to report pathway direction. It reports relative firing rate. At **fixed** wiring, lowering the receiving chamber's background drive moves asymmetry from 0.109 to 0.681; at **fixed** drive, raising the forward ratio

from 1 to 16 moves it only from 0.109 to 0.143. The measure's own range is [−1, 1], so this is not the whole of it — but drive moves the statistic several times further than the wiring it is supposed to be reporting. This is the most consequential single finding about the twin, and it propagates into every later section: any analysis that reads direction off coupling must control for rate first.

Three checks on the real devices confirmed the confound does not obviously bite there (firing-rate imbalance does not differ by device type), which is why §6.6 survives as a hypothesis.

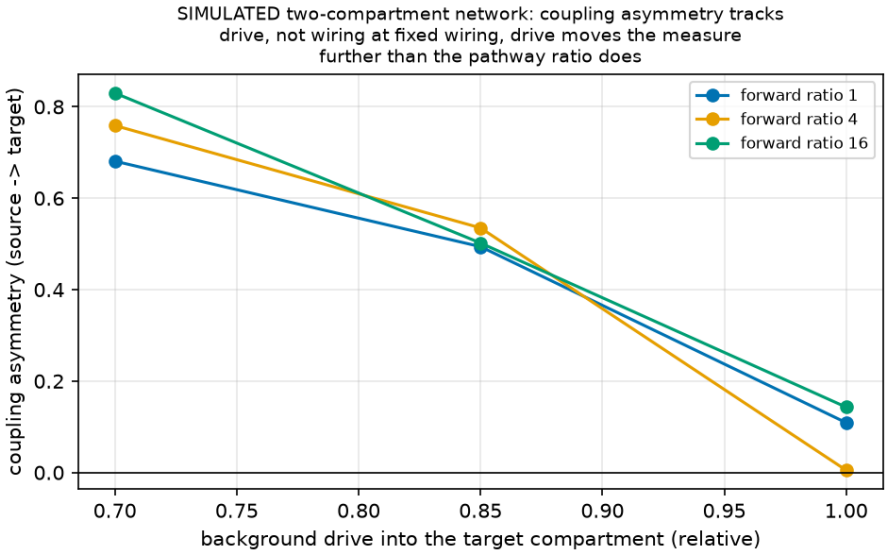


Figure 7: SIMULATED: measured coupling asymmetry against the forward/backward wiring ratio, at three levels of drive to the receiving chamber. Drive moves the measure further than wiring does.

7.3 Calibration

A random search over seven tuned parameters brought the simulated baseline into the real band for `mfr_hz`, `burst_rate_per_min` and `network_burst_rate_per_min`, and **not** for `pct_spikes_in_bursts` or `synchrony_index`. The residual mismatch is precisely the burst-structure and synchrony group that §6.5 and §6.6 independently identify as the unreliable part of the feature set across species. Three analyses agree about species; the platform contrast of §6.4 disagrees, and that disagreement is itself the finding.

7.4 Drug knobs

Each knob is compared against the real direction for its mechanism, scoring only features whose real direction is established by a cluster bootstrap:

knob	agrees / scored
GABA-A block	4 / 5
GABA-A potentiation	2 / 4
AMPA block	4 / 4
sodium-channel block	3 / 3

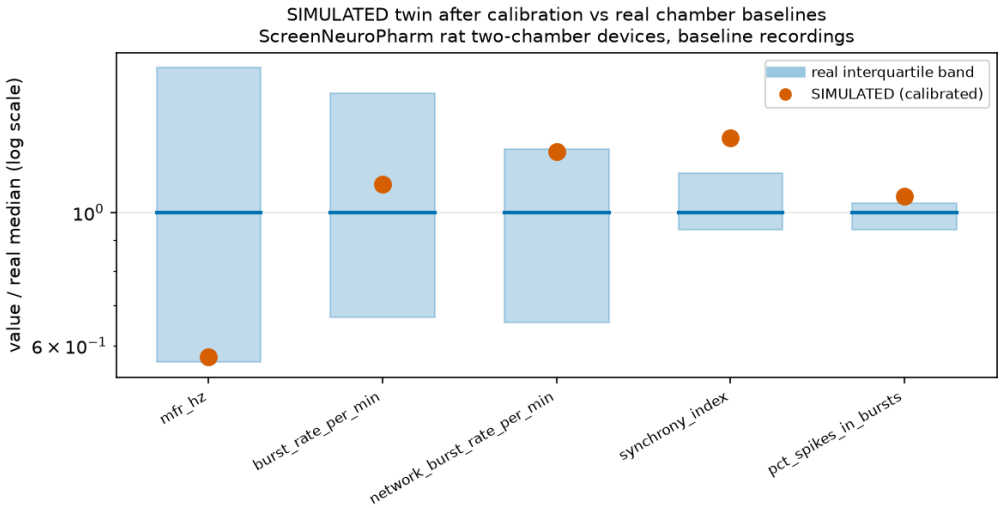


Figure 8: SIMULATED against REAL: calibrated baseline features of the twin compared with the real baseline band. Rate and burst-rate features land inside it; burst fraction and synchrony do not.

Directions only: each knob is compared against a real reference from a different platform, so this supports direction, never magnitude.

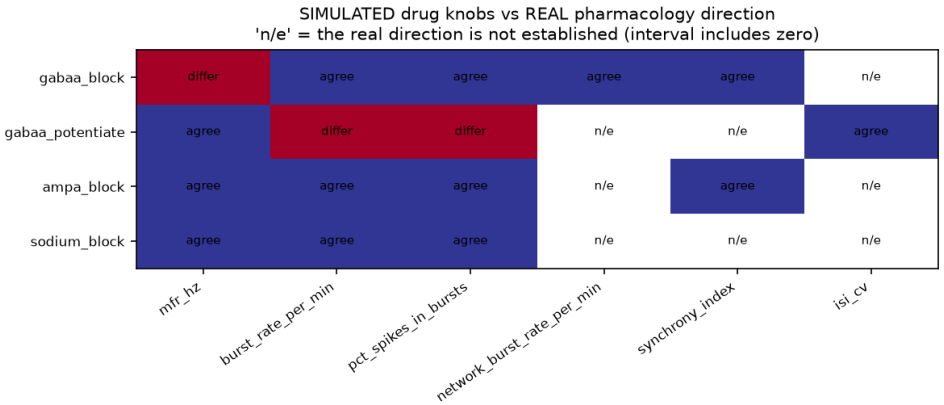


Figure 9: SIMULATED drug knobs against REAL pharmacology directions, per feature. Only features whose real direction is established by a cluster bootstrap are scored.

7.5 Simulation-based inference, and what it cannot do

A neural posterior estimator was trained on 3,500 simulations.

parameter	posterior contraction	KS statistic
`threshold_scale`	0.8	0.0571
`ampa_scale`	0.443	0.0871
`gabaa_scale`	0.323	0.1093
`inter_in_degree`	0.32	0.1029
`log10_forward_ratio`	0.311	0.0957

Only excitability is well identified, and only `threshold_scale` passes a rank-

uniformity check (p = 0.1956; every other parameter is below 0.01). **Posterior intervals from this estimator should not be quoted as calibrated.**

That statement is weaker than it looks, and the reason is worth recording. Posterior training was unseeded until this report was being built: rerunning identical code on the identical simulation cache moved contraction by about 0.02 and moved this p-value between 0.46, 0.0092 and 0.1956 — on one of those runs `threshold_scale` failed the check that it passes here. Training is now seeded and verified to reproduce across processes, but the run-to-run spread is the honest measure of how much weight a single rank-uniformity p-value can carry. Contraction was stable across the same runs; the rank test was not. A learning curve showed connectivity contraction still rising with training-set size but only slightly: more simulations alone will not fix it.

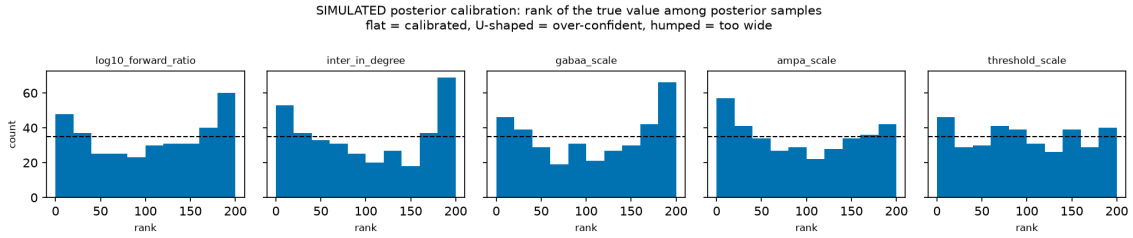


Figure 10: SIMULATED posterior calibration: rank of the true parameter among posterior samples. Flat is calibrated, U-shaped is over-confident. Only `threshold_scale` is close to flat.

Applied to 113 real two-chamber recordings from 16 devices, with 99.2% of real summary values inside the simulated range:

- The inferred GABA-A scale shows **no dose trend**: Spearman 0.118, 95% CI [−0.168, 0.419], clustered over devices.
- **Untreated recordings already infer 0.3781**, CI [0.3149, 0.4677], excluding the neutral value of 1.0. The knob is absorbing general model mismatch.
- Four summary dimensions — burst rate and active electrodes in both chambers — pass a range check but have almost no simulated draws where the real values sit. Range containment is not density containment.

An earlier version of this analysis produced an apparent dose trend (ρ = −0.197, p = 0.036). It was an artefact of pseudo-replication — 113 recordings from 16 devices treated as independent — and of posterior extrapolation on out-of-distribution dimensions. Both are now checks in code.

7.6 Directionality: what it takes to see asymmetry

120 simulated networks per arm with known wiring, symmetric against asymmetric. Electrode count is varied by subsampling one fixed array and recording length by truncating one simulation, so neither varies the culture.

The naive version is confounded exactly as §7.2 predicts: the full feature set separates the classes at AUC from 0.9219 to 0.9863, and three firing-rate features alone reach from 0.8644 to 0.9841.

Residualising every feature against both chambers' firing rates removes that, and carries its own validity check: rate features after residualising score from 0.3869

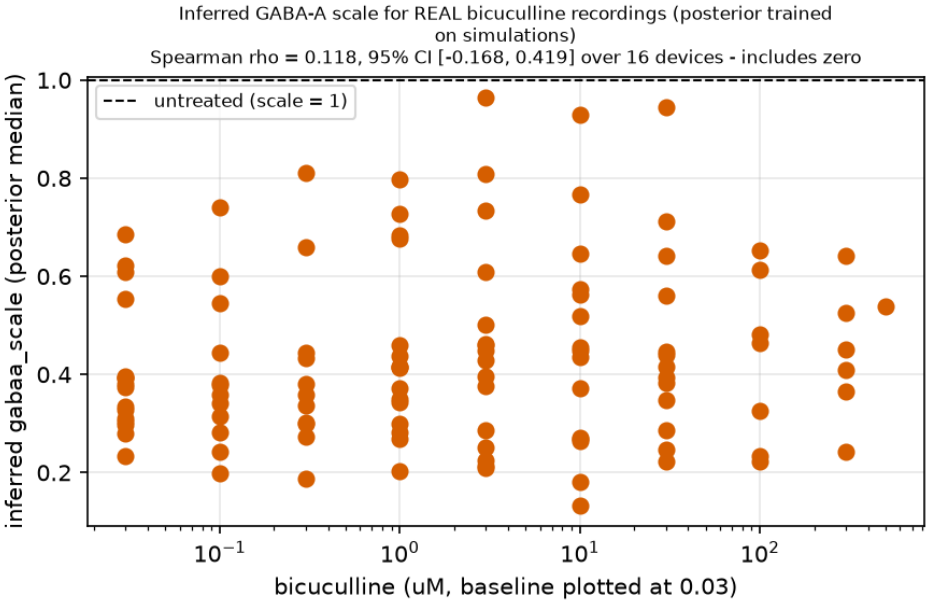


Figure 11: Inferred GABA-A scale for REAL bicuculline recordings, from a posterior trained on SIMULATIONS. There is no dose trend, and untreated chips already sit far below the neutral value of 1.

to 0.5969 and do not clear chance in every cell. Direction still separates at 0.9359, 95% CI [0.893, 0.9739], from a 10 s recording on 24 electrodes per compartment.

The rate-matched arm does not pass its own check, and is not the result. `metrics.json` states the criterion: in that arm, rate-only features should fall towards chance, and if they do not, the matching failed and the arm says nothing about direction. They do not — rate-only AUC there runs from 0.8587 to 0.9859. The drive adjustment did not remove the rate imbalance so much as reverse its sign, from −0.2756 and −0.5180 before to +0.3116 and +0.2377 after, leaving a worst relative gap of 0.26 against 0.2248 before it. That arm reaches AUC 0.9784 on only 6 electrodes, and an earlier draft of this report quoted exactly that number as the headline. It is withdrawn: the guardrail written into `metrics.json` says the arm is uninterpretable, and the correct response is to report the arm that passes, at the electrode count it actually needs.

Read this as a contrast, not a measurement. The same summaries that classify two well-spaced wiring regimes at 0.98 give a posterior contraction of only 0.311 for the continuous ratio (§7.5). Classifying is a far weaker demand than estimating.

7.7 The same question on real devices, and a null

§7.6 could only ask its question in simulation, where the ground truth is something we set. The Mateus deposit (§4.2) has it fabricated into the device: straight `control` microchannels against `arrows`, `rams` and `tesla`-valve geometries, 57 recordings from 26 chips.

The deposit publishes spike times and the measurement but not the electrode map it was computed on, so the geometry is recovered with the authors' own criterion — an event on all five electrodes of a microchannel with at most 1 ms between consecutive electrodes. That finds 10–15 channels per chip in a single row band,

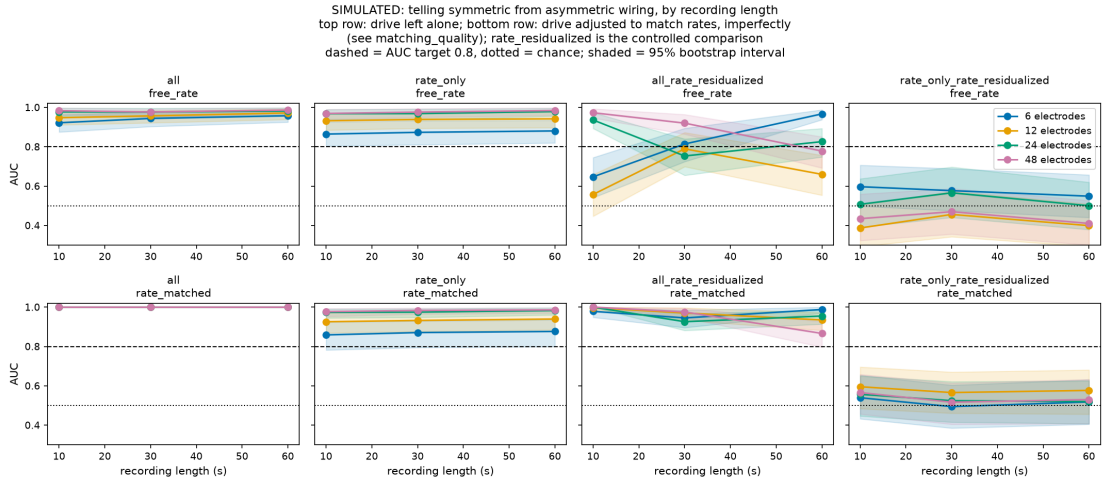


Figure 12: SIMULATED: telling symmetric from asymmetric wiring, by recording length and electrode count. Top row leaves drive alone; the bottom row adjusts drive to match rates and does so imperfectly, which is why that arm is reported as inconclusive. The rate-residualised columns are the controlled comparison, and the rightmost column is the validity check itself.

which is what the published five-electrode design predicts. Direction is reported as a magnitude: the deposit does not record how each device was seated on the array, so |asymmetry| says how one-sided propagation is, not which way it points.

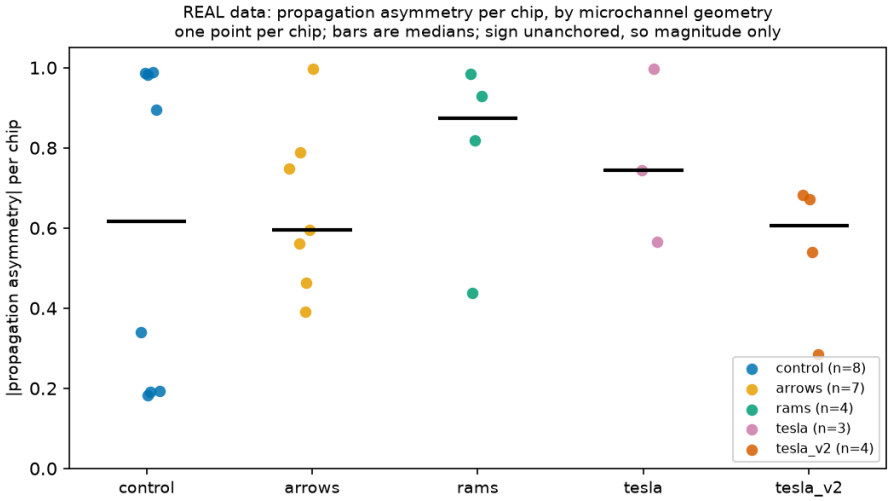


Figure 13: REAL data: propagation asymmetry per chip by microchannel geometry. One point per chip, bars are medians. The control group spans almost the whole range on its own.

The unit decides the answer. A chip recorded at DIV 12, 17 and 21 supplies three recordings and one device. Counting recordings, median |asymmetry| runs from 0.339 for control to 0.854 for rams, which looks like a geometry effect. Counting chips, it is 0.617 against 0.875 with Kruskal $p = 0.64$ and every Holm-adjusted comparison against control at 1.0. Both are kept in `metrics.json`, the per-recording one labelled as the wrong unit, because the gap between them is the result. This is the pseudo-replication of §7.5 a third time, on a different dataset.

The statistic is also confounded with sampling. |asymmetry| is a ratio of

counts, so a chip with few propagation events looks one-sided by chance; across chips the correlation between event count and |asymmetry| is −0.494, p = 0.010. Restricting to the better-sampled half leaves 13 chips and p = 0.086, which is not a result either.

What this does not say. It does not contradict the deposit's own paper: those authors had the electrode map, measured per-channel propagation velocities, and asked a different question. The claim here is narrower — this statistic, aggregated per chip, does not separate these geometries, on 26 chips in five groups of three to eight, with control chips confined to experiments 9–11 so that condition is confounded with culture batch by the deposit's design.

What it says about the twin. §7.6's directionality result is simulation-only, and this is the first time that question has been put to devices whose asymmetry was fabricated rather than simulated. It does not reproduce. The sample is too small and too confounded to call that a refutation, and too directly aimed to leave out.

---

8. Limitations

Data. The asset is 21,099 rat records against 173 human; every human claim rests on two sources. Most records are curve fits from which our features cannot be recomputed. Dropping Tukker 2020 on license grounds leaves the transfer study on a single source. Compound coverage is thin: six mechanism classes, 73 pairs.

Models. Balanced accuracies around 0.5 on six classes are above chance but far from usable; intervals are wide and overlapping. Conformal sets cover honestly but are too large to act on. Calibration error may be as high as 0.2557. The nine-feature vector is effectively a rate feature plus noise for this task. The implementation uses single-level grouped CV where the plan specified nested CV, so hyperparameter choice is not fully held out.

The twin. Coupling asymmetry tracks rate, not direction (§7.2). Calibration misses burst structure and synchrony (§7.3). Connectivity parameters are weakly identified and no parameter's posterior is calibrated except by the weakest test (§7.5). On real data the drug knob absorbs model mismatch (§7.5). Simulated recordings are 10 s against real recordings of 480 s. The simulated array is one fixed geometry while real chambers carry 48, 36 or 60 electrodes. HD16's positive result is identifiability under a correct model, since the same simulator generated both classes, and its rate controls remove only linear rate dependence.

The real-device test. The §7.7 null rests on 26 chips in five groups, with condition confounded with culture batch by the deposit's design, an unanchored direction, and a statistic that varies with how many propagation events a chip produced. It bounds the effect rather than excluding it. The microchannel geometry is also recovered from the spike times rather than read from a map the deposit does not publish, so it approximates the authors' own analysis rather than reproducing it.

Process. A CI workflow is configured and the work described here has been pushed to it, but no CI result has been observed from the environment these results were produced in; treat CI as unverified. The demo video specified for Phase 4 has not been recorded. Every twin metric produced before the coupling-definition fix is incomparable with those after it, and was regenerated.

A fuller list, with dates and commits, is in `report/LAB_NOTEBOOK.md` (33 numbered limitations).

---

9. Reproducibility

```
pip install uv
uv venv --python 3.11 .venv
uv pip install --python .venv/bin/python -e ".[dev]"
python run_all.py --quick     # every figure and metric, reduced settings
python run_all.py             # full run
```

- `--quick` completes in **401 s on a fresh clone** of this repository — install, download, build, model, twin and demo — against a 600 s budget, and each run records its own per-stage timings under `metrics.run_all_timings_quick`. Most of the variable part is downloading: the same clone measured 52 s and 133 s for the fetch stage on two runs; the full twin stage takes roughly 2,000 s and records `metrics.run_all_timings_full`. Both are logged under `outputs/logs/` as well.
- **`--quick` is a path test, not a reproduction of the report.** Measured on a fresh clone, it fetches the small-tier sources only and therefore builds a much smaller asset than the one this report describes; it produces 8 of the 14 figures and omits six `metrics.json` entries — `toxcast_mea_inventory`, `twin_hd12_calibration`, `twin_hd15_real_inference`, `compartment_asymmetry`, `microchannel_directionality` and the full-run timings. The analyses behind those either need a full-tier download (the 109 MB microchannel archive, the ToxCast release) or need compartment records the small-tier asset does not contain, and each skips itself with a logged reason rather than failing. The §7.6 sweep additionally has too few simulations per cell to fit a discriminator, so its `cheapest_configuration` entries are written as `null`.
- **Reproducing the numbers in this report takes the full path**: `python run_all.py` after a full-tier fetch. `metrics.json` records which mode produced each stored value, and quick and full runs write to separate roots (`outputs/` against `outputs/quick/`) so neither can overwrite the other.
- Dependencies are pinned in `pyproject.toml`; numpy is held at 2.3.5 because brian2 2.9.0 does not import against 2.4.
- Every number in this report is produced by a script under `src/` and written to `outputs/metrics.json`; `python -m neuromps.report.verify` checks the report against it.
- Caches refuse to be reused across a change in what they mean: the simulation cache stores its summary names, electrode count and a hash of the source of every function that defines a summary.
- Tests run with `pytest` (feature extraction, QC flags, provenance hashing, the MEA simulator, and a contract test that fails if any headline metric loses its confidence interval).

9.1 Where the time goes

Each run writes its own stage timings, so this is generated rather than transcribed: see `metrics.run_all_timings_quick.by_stage_s` and `run_all_timings_full`. On

the fresh-clone run above, the twin stage dominates, then downloading, then the asset build.

9.2 What the repository enforces, rather than documents

- **Provenance.** The asset build fails if any record's provenance hash disagrees with the manifest.
- **One feature extractor.** The twin computes features through the same `features/mea.py` as real recordings, and — after a defect that survived for several phases — through the same coupling measure on the same per-electrode input.
- **Interval discipline.** A contract test fails if any headline metric loses its confidence interval, and the resample count is recorded next to every estimate.
- **Cache meaning.** The simulation cache stores its summary names, its electrode count and a hash of the source of every function that defines a summary; it refuses to load across a change in any of them.
- **Report numbers.** `neuromps.report.verify` re-reads every citation in this document and fails on disagreement.
- **No clinical claims.** A test scans the reports for efficacy, diagnostic and treatment language and fails unless the line is a disclaimer ruling it out.

---

10. The demo

`streamlit run app/app.py`, after `python run_all.py --quick`, opens three tabs — the data asset and its standard, mechanism prediction with a conformal set beside the transfer results, and the twin with its posterior and what it does on real recordings.

No interaction waits on a simulator: `run_all.py --stage demo` precomputes a grid of 75 twin simulations spanning the prior with posterior samples per point, posterior samples for all 113 real two-chamber recordings, and a classifier fitted once, in 85.1 s. Through Streamlit's own test harness every interaction lands between 0.6 and 0.9 s; the exception is a button running one Brian2 simulation at the current slider values in about 4.5 s, so the twin is verifiably a simulator rather than a lookup table. The app reads `outputs/metrics.json` and `outputs/demo/` and recomputes nothing, so it cannot disagree with this report. Tab-by-tab detail is in `app/README.md`; screenshots are in `outputs/figures/app_tab*.png` and in the write-up.

---

11. AI assistance disclosure

This work was produced in a paired workflow with Anthropic's Claude (Claude Opus 5) via Claude Code. The model wrote the large majority of the code, tests and prose in this repository, under step-by-step human direction: the human set the plan and phase gates, chose the scope, approved or rejected each step, and directed the scientific decisions.

Specific practices used to keep that honest, all visible in the repository history:

- Numbers are never written by hand into documents. Every figure in this report is emitted by a script into `outputs/metrics.json` and re-checked against it by `neuromps.report.verify`, which caught a stale timestamp and three misanchored citations while this report was being written.
- Results were reported as observed, including when they contradicted the plan's expectation — the negative findings in §6.2, §6.3, §7.2 and §7.5 are examples.
- `report/LAB_NOTEBOOK.md` records what was tried and abandoned, with dates and commit hashes, including several defects the model introduced and then found in its own work.

---

12. Sources and licenses

Data sources, with licenses and citations, are listed in `CITATIONS.md` and recorded per file in `data/raw/MANIFEST.json`. Third-party software licenses are in `THIRD_PARTY.md`. This repository is released under the MIT license (`LICENSE`).

Appendix A. Where each result lives

Every number in this report comes from a key in `outputs/metrics.json`, written by a script under `src/`. The main entry points:

result	metrics key	script
asset counts and provenance	`data_asset`, `data_card`	`neuromps.etl.build`
mechanism classification, calibration	`mechanism_baselines`	`neuromps.models.mechanism`
conformal sets, ablations	`uncertainty_and_ablations`	`neuromps.models.conformal`, `.ablation`
species transfer and shift	`species_transfer`	`neuromps.models.transfer`
compartment asymmetry	`compartment_asymmetry`	`neuromps.models.asymmetry`
twin baseline and drive sweep	`twin_hd11_validation`	`neuromps.twin.validate`
twin calibration	`twin_hd12_calibration`	`neuromps.twin.calibrate`
drug knob directions	`twin_hd13_perturbation`	`neuromps.twin.perturbation`
posterior training and coverage	`twin_hd14_inference`	`neuromps.twin.inference`
posterior applied to real recordings	`twin_hd15_real_inference`	`neuromps.twin.real_inference`
directionality experiment	`twin_hd16_directionality`	`neuromps.twin.directionality`

Figures are written to `outputs/figures/` (real data) and `outputs/twin/` (twin output), the latter carrying `simulated` in the filename except for `hd15_real_dose_response.png`, which plots real recordings under a simulation-trained model and says so in its caption. The rule is enforced in the captions, not the filenames: no simulated output is presented as measurement.